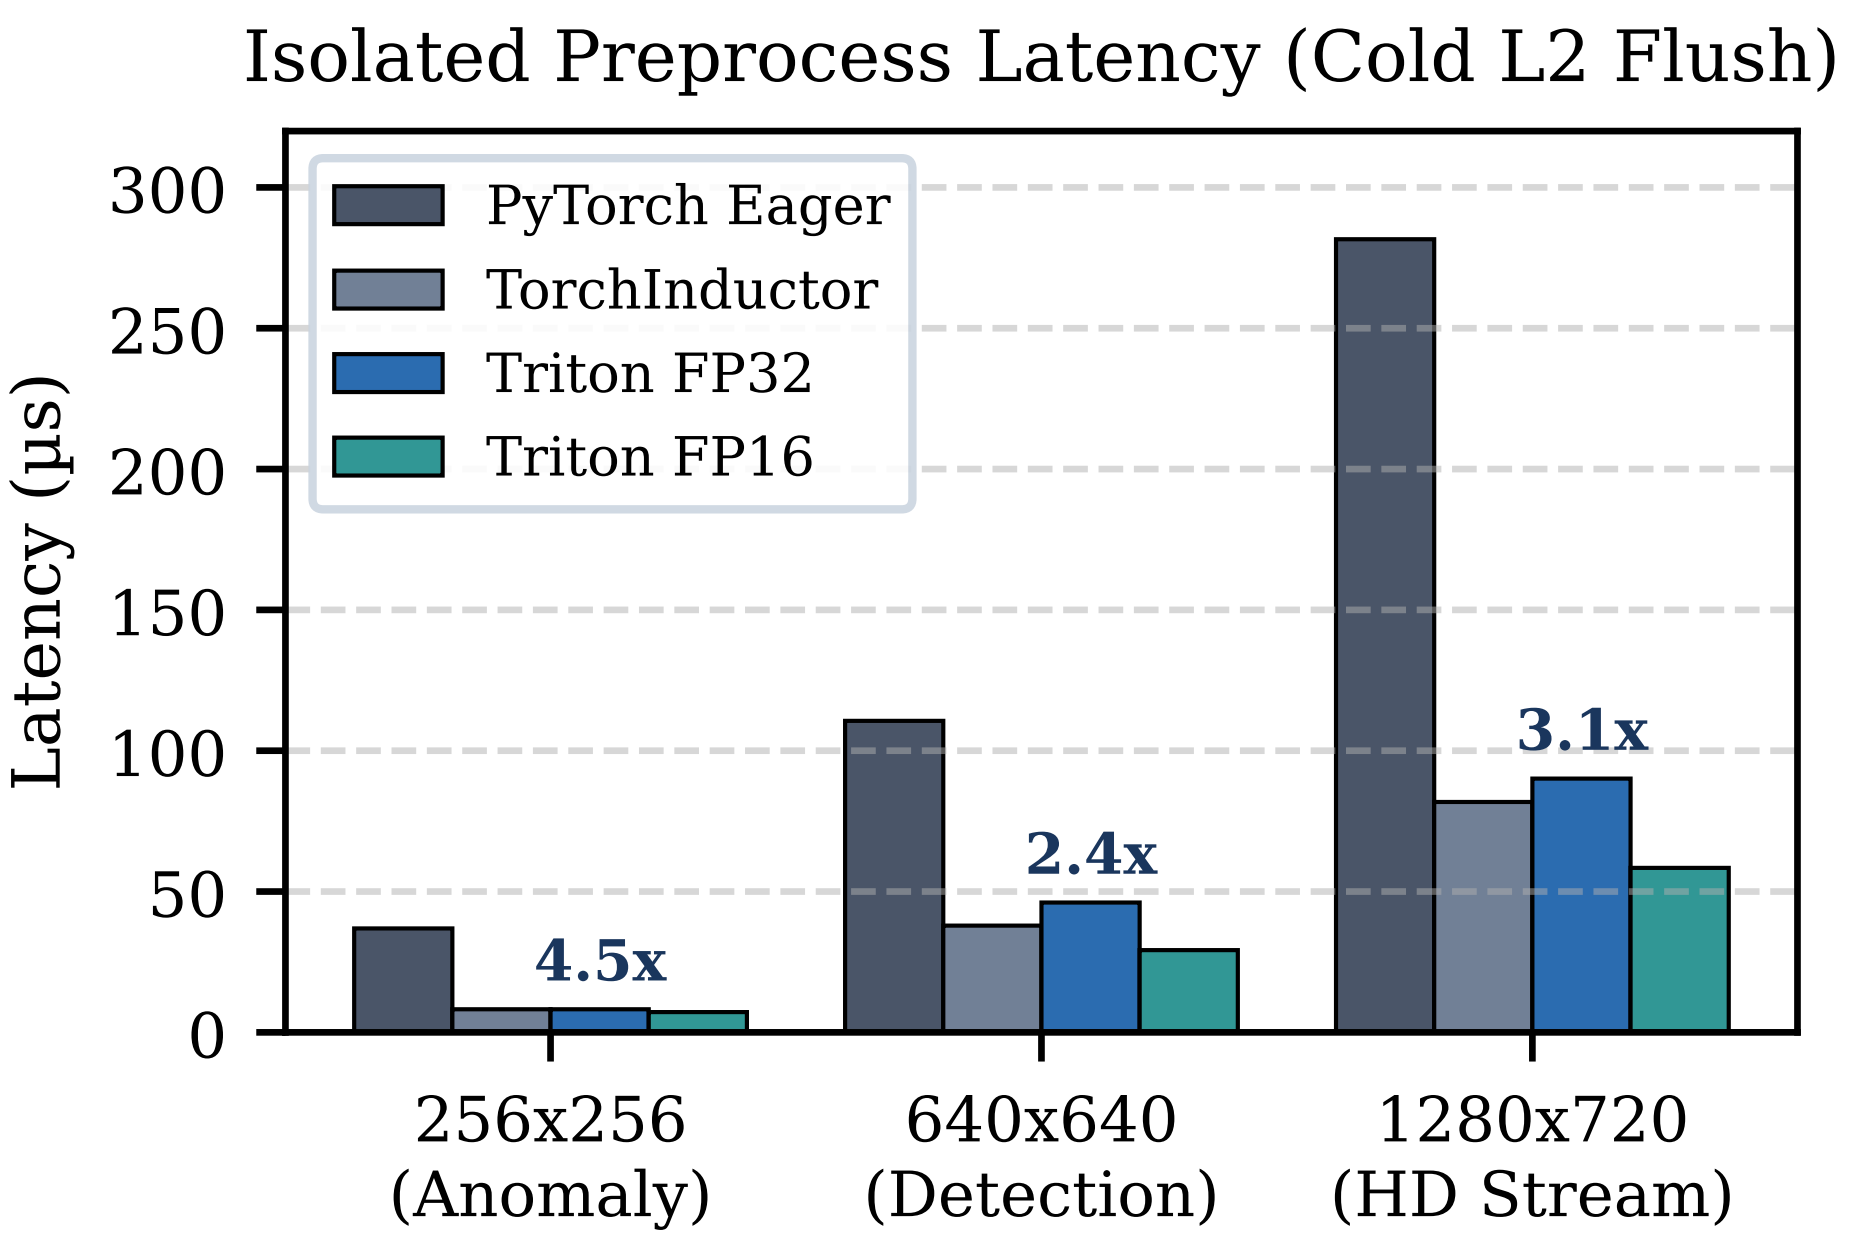
Isolated Preprocess Latency (Cold L2 Flush)
PyTorch Eager
TorchInductor
Triton FP32
Triton FP16
Latency (µs)
300
250
200
150
100
50
0
4.5x
2.4x
3.1x
256x256
(Anomaly)
640x640
(Detection)
1280x720
(HD Stream)